Keith Tang
Kiran Khatri

4. We believe that an unordered_map is the most efficient data structure for Amazon340, as the purpose of Amazon340 is to manage vendor's products.

For data volume, a vendor can have a wide range of products, it can range from a few to thousands or more. Therefore, the space complexity of the data structure has to be low to compatible with a large number of products.

For operation frequency, Amazon340 will have a high frequency of searches. Many customers will be doing multiple searches to retrieve product information. Therefore, the time complexity of search operation has to be low to ensure a fast retrieval of products.

For data growth, it is expected to grow gradually, as it is less likely that the vendor will add a large number of new type of products. When the products are being added overtime, we do not expect an exponential growth. Therefore, the time complexity of insert operation does not need to be low, but it should still be efficient as possible.

For search needs, the queries may be complex in reality to search for products, but it is simple in Amazon340. There is very less information for products in Amazon340, so it is enough for the query to support searches by key data (product name) only.

An unordered_map has O(1) average time complexity and O(n) at worst case for insert, delete and search operations, it will be the most efficient before the hash collides occur, which make the worst case happens. Also, the space overhead is the pointers in hash table, it is not the smallest but is generally small compared to other data structures.